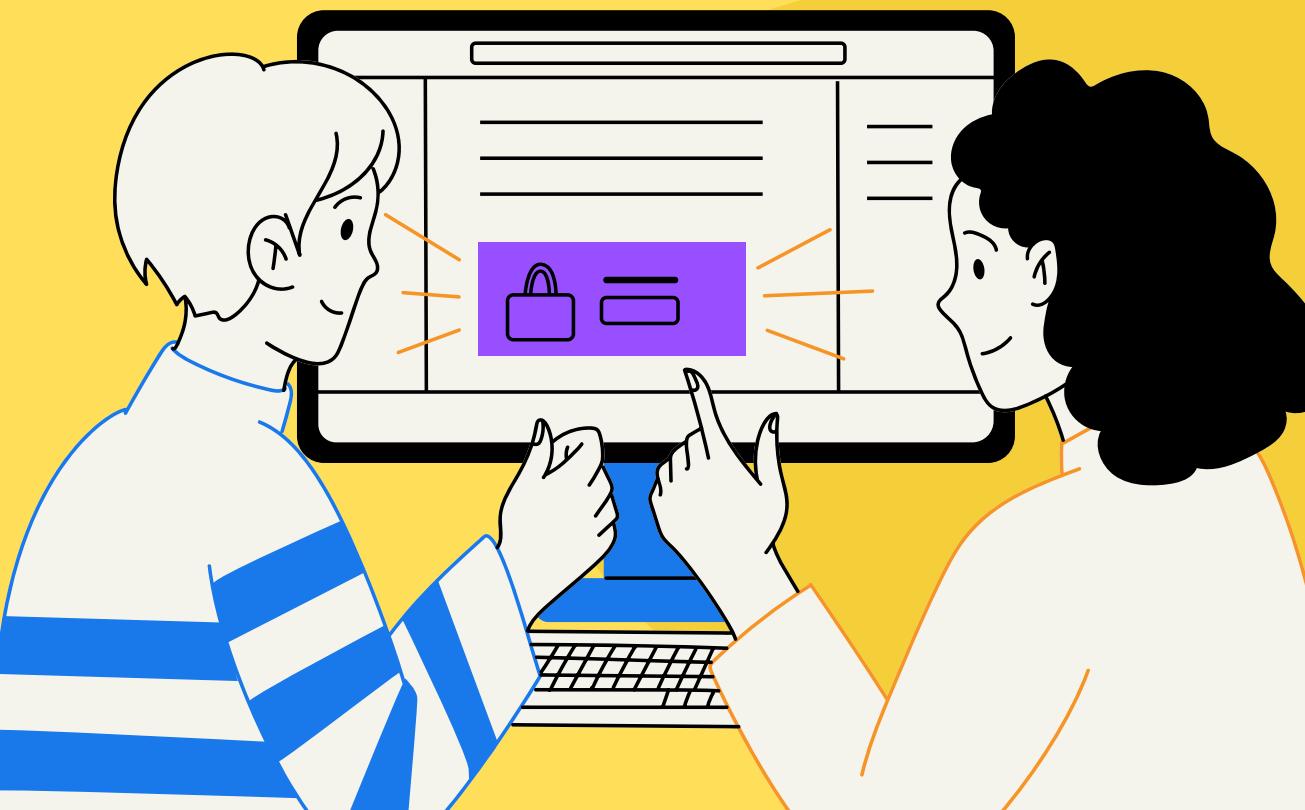


# Module 6

# ADVANCED CONCURRENCY AND SYNCHRONIZATION



# SYNCHRONIZATION TECHNIQUES



# `sync.WaitGroup`

Used to wait for a group of goroutines to finish executing. Helps coordinate concurrent tasks.

# `sync.Mutex`

A mutual exclusion lock used to protect shared resources from concurrent access.





# sync.RWMutex

A read/write mutex that allows multiple readers but only one writer at a time.

## Key Points to Remember

- Deadlocks occur when goroutines wait indefinitely due to cyclic dependencies on locks.
- Race Conditions happen when multiple goroutines access and modify shared resources without synchronization.
- Always unlock a sync.Mutex or sync.RWMutex after locking to avoid deadlocks (defer m.Unlock() is a best practice).

# WORKER POOLS AND PIPELINES



# Worker Pool

A pattern where a fixed number of goroutines process a large number of tasks from a queue (channel).



# Pipeline

A chain of stages where data flows from one stage to another via channels.



# Key Points to Remember

- **Avoid excessive goroutines:** Spawning too many goroutines without management can lead to high memory usage and scheduling overhead.
- **Use buffered channels:** Helps manage task queue size and prevents blocking.
- **Graceful shutdown:** Always close channels properly to prevent goroutines from getting stuck.

# CONTEXT PACKAGE



# `context.Background()`

The root context, used as a base for other contexts.

# `context.WithCancel()`

Creates a child context that can be canceled explicitly.

# `context.WithDeadline()`

Sets an absolute time limit after which the context expires.



# Key Points to Remember

- Use `context.WithCancel()` to stop goroutines gracefully:  
**Always pass ctx down to goroutines and check ctx.Done().**
- **Avoid goroutine leaks:** Ensure that goroutines properly exit when the context is canceled.
- Set timeouts for network calls to prevent hanging operations (e.g., `context.WithTimeout()` for HTTP requests).

# REAL-WORLD CONCURRENCY PATTERNS



# Producer-Consumer

- Producers generate data; consumers process it.
- Uses a queue (buffer) for decoupling.
- Requires synchronization (locks, semaphores).
- Used in job scheduling, logging, and message queues.



# Fan-In

- Multiple producers → One consumer (merging data streams).
- Reduces contention, aggregates results efficiently.

Used in data processing, event aggregation.



# Fan-Out

- One producer → Multiple consumers (distributing tasks).
- Improves parallel processing and system scalability.
- Used in load balancing, task distribution (e.g., web servers).



# THANK YOU

